



Stable SSAO in Battlefield 3 with Selective Temporal Filtering

Louis Bavoil

Developer Technology Engineer
NVIDIA

Johan Andersson

Rendering Architect
DICE / EA

SSAO

- Screen Space Ambient Occlusion (SSAO)
 - Has become de-facto approach for rendering AO in games with no precomputation
 - Key Idea: use depth buffer as approximation of the opaque geometry of the scene
 - Large variety of SSAO algorithms, all taking as input the scene depth buffer

[Kajalin 09]
[Loos and Sloan 10]
[McGuire et al. 11]

HBAO

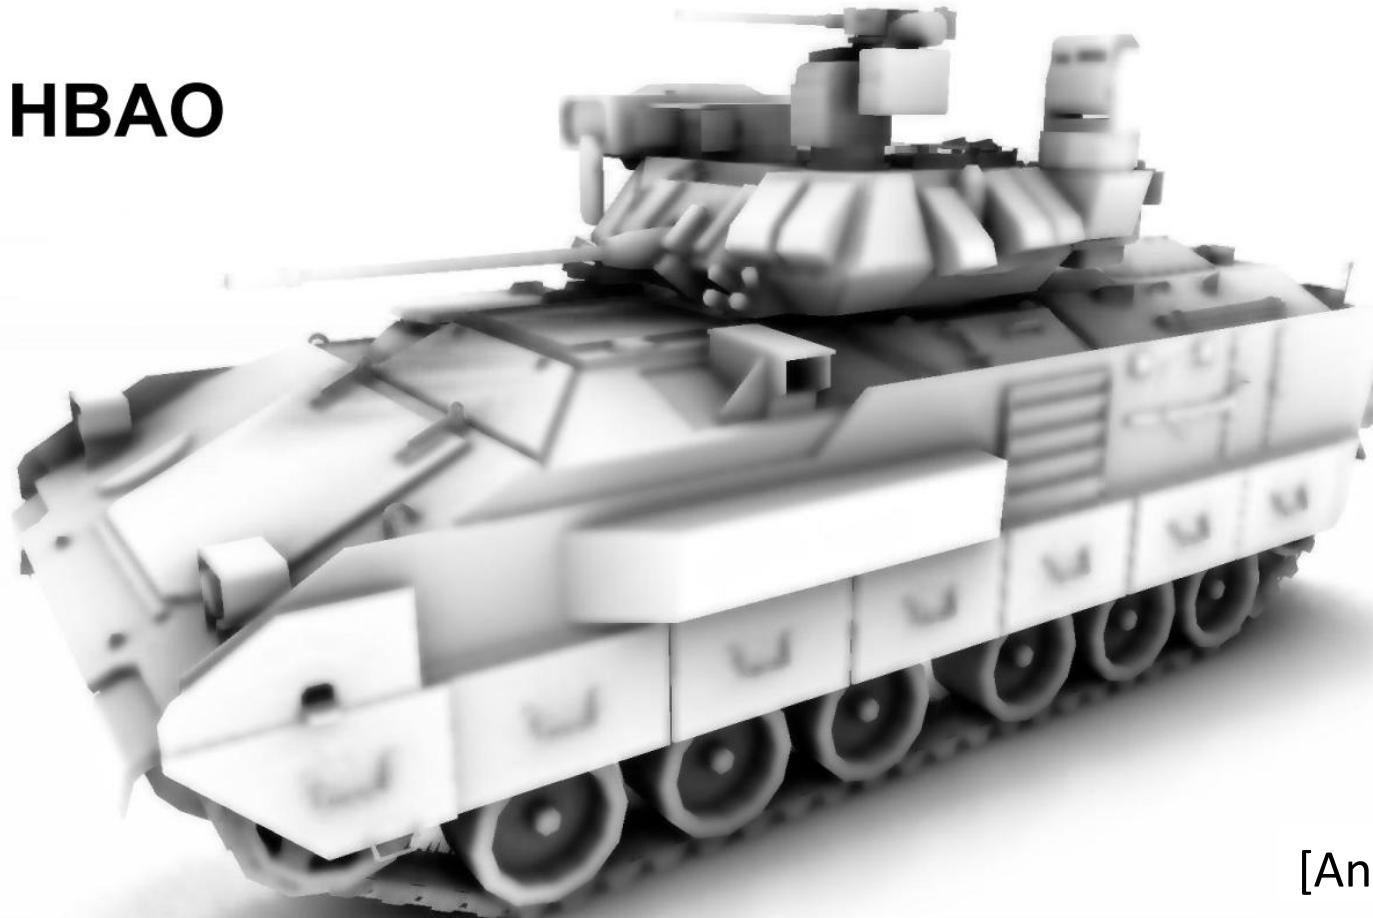
- Horizon-Based Ambient Occlusion (HBAO)
 - Considers the depth buffer as a heightfield, and approximates **ray-tracing** this heightfield
 - Improved implementation available in NVIDIA's SDK11 (SSAO11.zip / HBAO_PS.hlsl)
 - Used for rendering SSAO in Battlefield 3 / PC for its “High” and “Ultra” Graphics Quality presets

[Bavoil and Sainz 09a]

[Andersson 10]

[White and Barré-Brisebois 11]

HBAO



[Andersson 10]

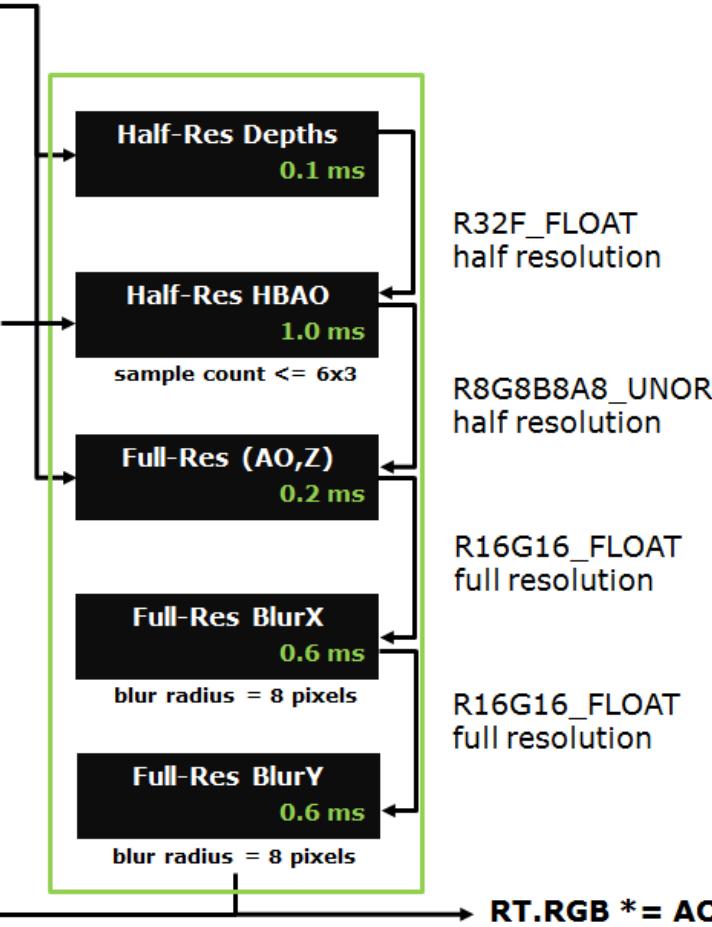
View-Space Depths

R32F_FLOAT
full resolution

4x4 Jitter Texture

R16G16B16A16_SNORM
frame independent

RT = Shaded Colors
R16G16B16A16_FLOAT
full resolution



Original HBAO Implementation in Frostbite 2

Frame time (GPU):

25.2 ms

Total HBAO time:

2.5 ms (10% of frame)

[1920x1200 “High” DX11
GeForce GTX 560 Ti]

Screenshot: With HBAO



Screenshot: HBAO Only

BETA



The HBAO looked good-enough on screenshots...

Video: Flickering HBAO

BETA



...but produced objectionable flickering in motion on thin objects such as alpha-tested foliage (grass and trees in particular)

Our Constraints

- PC-only
 - In Frostbite 2, HBAO implemented only for DX10 & DX11
- Low Perf Hit, High Quality
 - Whole HBAO was already 2.5 ms (1920x1200 / GTX 560 Ti)
 - HBAO used in High and Ultra presets, had to look great

Considered Workarounds

- Full-resolution SSAO or dual-resolution SSAO (*)
...but that more-than-doubled the cost of the SSAO, and some flickering could remain
- Brighten SSAO on the problematic objects
...but we wanted a way to keep full-strength SSAO on everything (in particular on foliage)

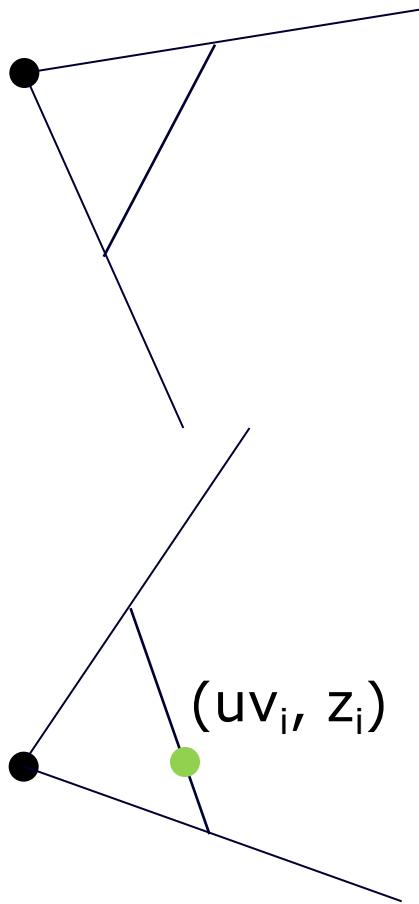
(*) [Bavoil and Sainz 09b]

Temporal Filtering Approach

- By definition, AO depends only on the scene geometry, not on the camera
 - For static (or nearly-static geometry), can re-project AO from previous frame
 - Reduce AO variance between frames by using a temporal filter: $\text{newAO} = \text{lerp}(\text{newAO}, \text{previousAO}, x)$
- Known approach used in Gears of War 2

[Nehab et al. 07]
[Smedberg and Wright 09]

Previous
Camera
(Frame $i-1$)

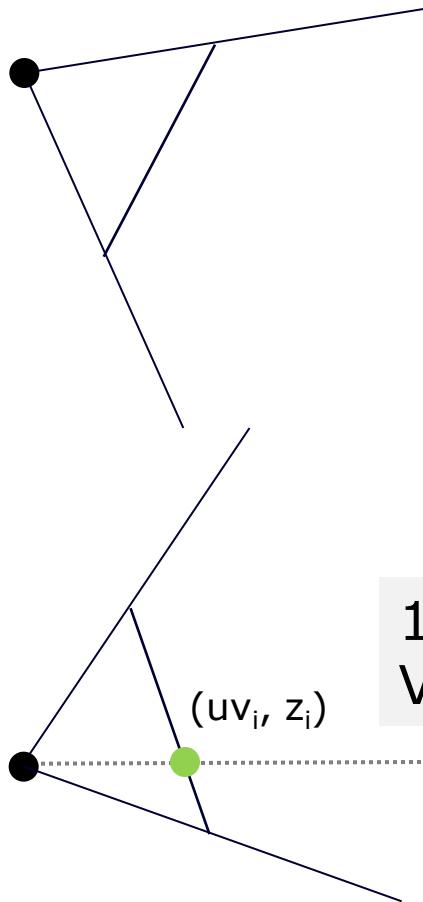


Current
Camera
(Frame i)

Reverse Reprojection

[Nehab et al. 07]

Previous
Camera
(Frame $i-1$)



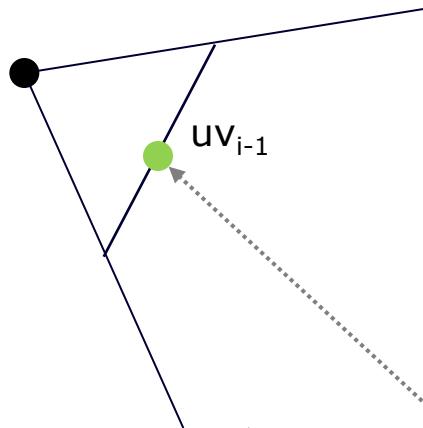
Reverse Reprojection

[Nehab et al. 07]

1. Current
ViewProjection $^{-1}$

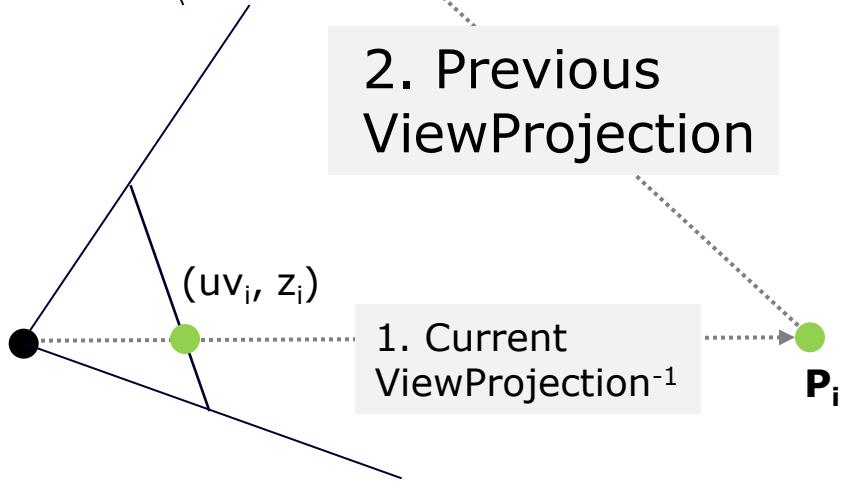
P_i

Previous
Camera
(Frame $i-1$)



2. Previous
ViewProjection

Current
Camera
(Frame i)

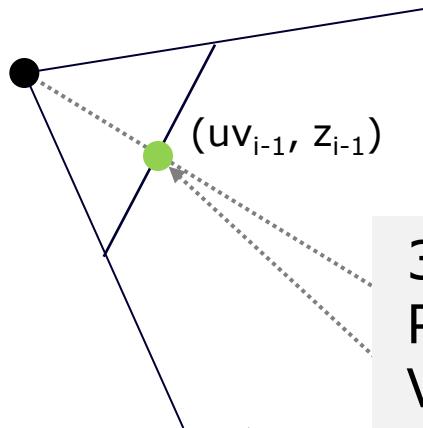


1. Current
ViewProjection $^{-1}$

Reverse Reprojection

[Nehab et al. 07]

Previous
Camera
(Frame $i-1$)



3. Fetch
Previous
View Depth

Current
Camera
(Frame i)

2. Previous
ViewProjection

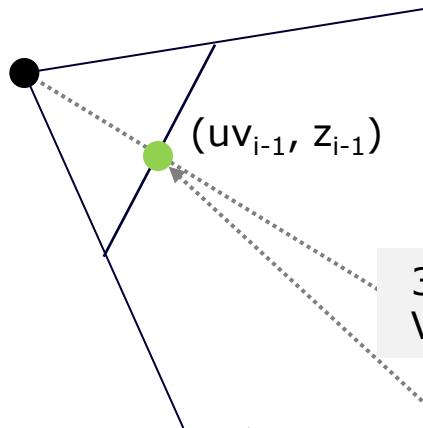
1. Current
ViewProjection $^{-1}$

Reverse Reprojection

[Nehab et al. 07]

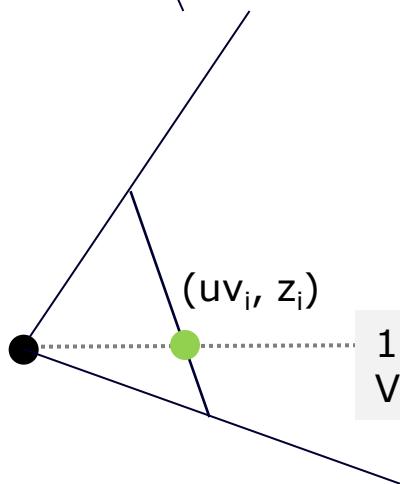
P_{i-1}
 P_i

Previous
Camera
(Frame $i-1$)



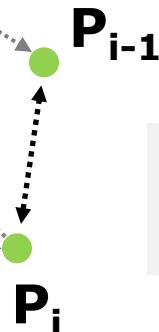
3. Fetch Previous
View Depth

Current
Camera
(Frame i)



1. Current
ViewProjection $^{-1}$

2. Previous
ViewProjection

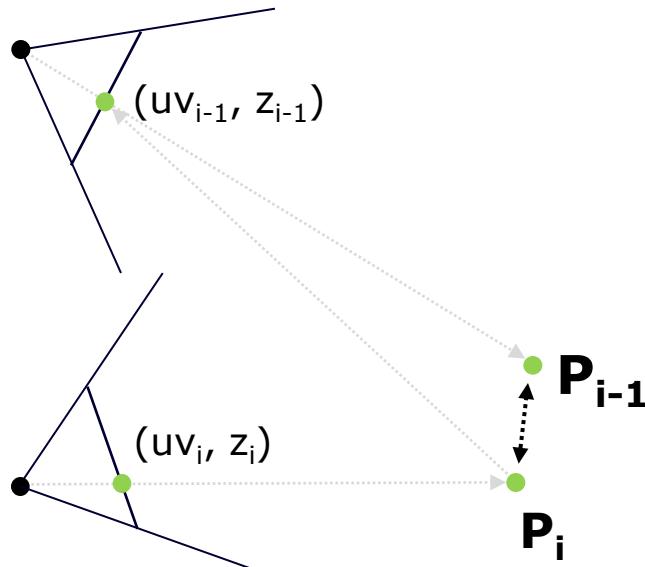


4. If $P_i \sim= P_{i-1}$
re-use AO(P_{i-1})

Reverse Reprojection

[Nehab et al. 07]

Temporal Refinement [Mattausch et al. 11]



If $P_{i-1} \sim = P_i$

$$AO_i = (\mathbf{N}_{i-1} AO_{i-1} + AO_i) / (\mathbf{N}_{i-1} + 1)$$

$$N_i = \min(\mathbf{N}_{i-1} + 1, N_{\max})$$

Else

$$AO_i = AO_i$$

$$N_i = 1$$

N_i = num. frames that have been accumulated in current solution at P_i

N_{\max} = max num. frames (~ 8), to keep AO_{i-1} contributing to AO_i

Disocclusion Test [Mattausch et al. 11]

$$P_i \approx P_{i-1} \Leftrightarrow \left| 1 - \frac{w_i}{w_{i-1}} \right| < \varepsilon$$

w_i = ViewDepth(View_{i-1}, P_i)

w_{i-1} = ViewDepth(View_{i-1}, P_{i-1})

Relaxed Disocclusion Test

$$P_i \approx P_{i-1} \Leftrightarrow \left| 1 - \frac{w_i}{w_{i-1}} \right| < \varepsilon$$

To support nearly-static objects

- Such as foliage waving in the wind (grass, trees, ...)
- We simply relaxed the threshold (used $\varepsilon = 10\%$)

Video: Temporal Filtering with $\epsilon=+\text{Inf}$



Flickering is fixed, but there are ghosting artifacts due to disocclusions

Video: Temporal Filtering with $\epsilon=0.1\%$

77



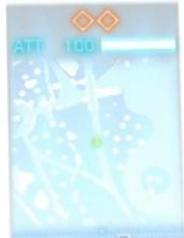
Flickering on the grass (nearly static), but no ghosting artifacts

Video: Temporal Filtering with $\epsilon=10\%$



No flickering, no ghosting, 1% perf hit (25.2 -> 25.5 ms)

Video: Trailing Artifacts



New issue: trailing artifacts on static objects receiving AO from dynamic objects

Observations

1. With temporal filtering OFF

The **flickering pixels** are mostly on foliage. The other pixels do not have any objectionable flickering.



2. With temporal filtering ON

The **trailing artifacts** (near the character's feet) are not an issue on foliage pixels.

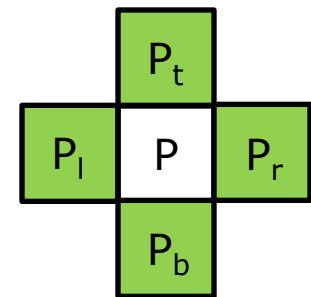


Selective Temporal Filtering

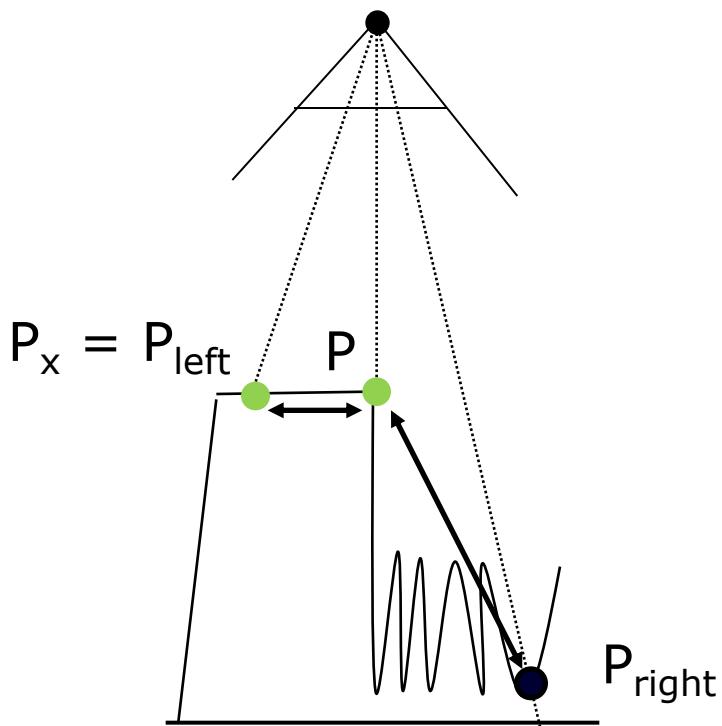
- Assumption
 - The set of **flickering pixels** and the set of **trailing pixels** are **mutually exclusive**
- Our Approach:
 1. Classify the pixels as **stable** (potential trailing) or **unstable** (potential flickering)
 2. Disable the temporal filter for the stable pixels

Pixel Classification Approach

- Normal reconstruction in SSAO shader
 - $P_x = ||P - P_{left}|| < ||P - P_{right}|| ? P_{left} : P_{right}$
 - $P_y = ||P - P_{top}|| < ||P - P_{bottom}|| ? P_{top} : P_{bottom}$
 - $N = \pm \text{normalize}(\text{cross}(P - P_x, P - P_y))$
- Idea: If reconstructed normal is noisy, the SSAO will be noisy

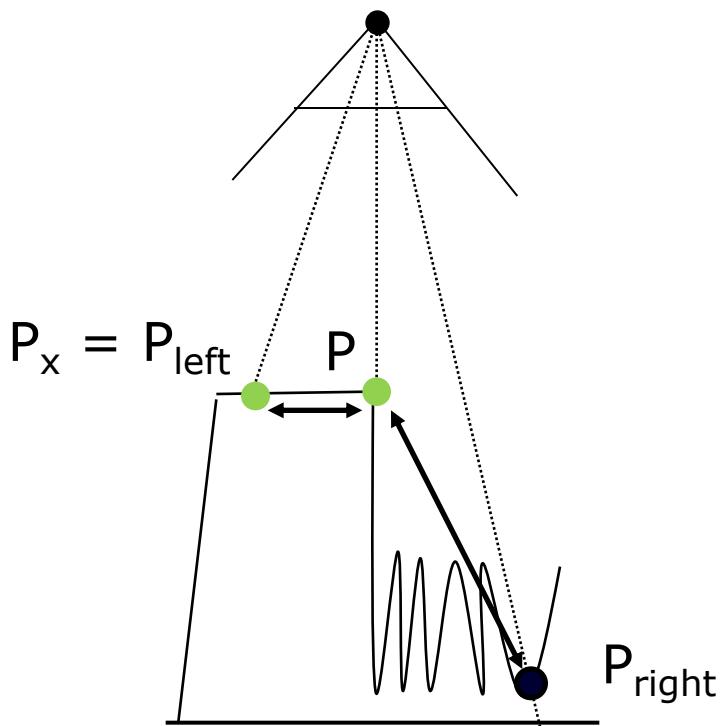


Piecewise Continuity Test



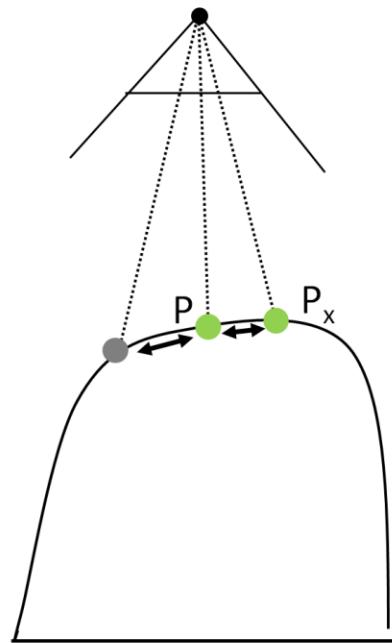
1. Select nearest neighbor P_x between P_{left} and P_{right}

Piecewise Continuity Test

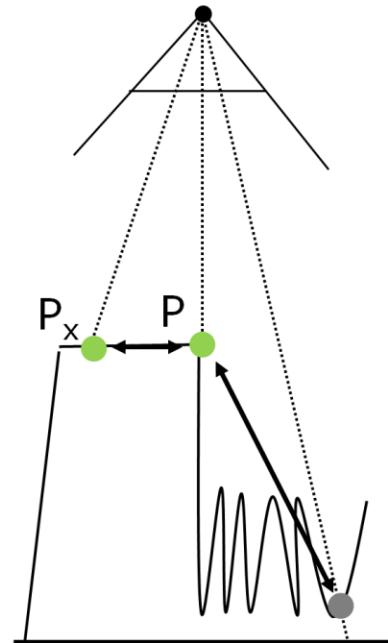


1. Select nearest neighbor P_x between P_{left} and P_{right}
2. Continuous pixels \Leftrightarrow
$$|| P_x - P || < L$$
where L = distance threshold
(in view-space units)

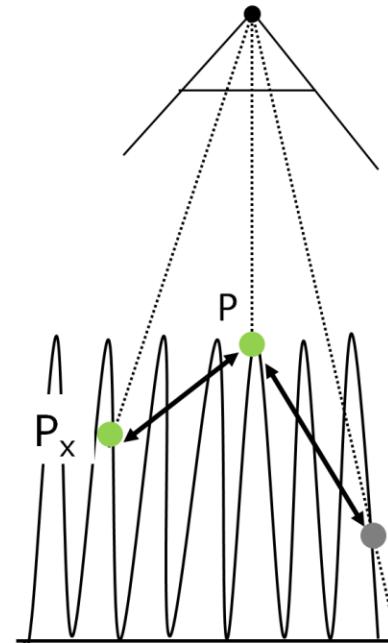
Pixel Classification Examples



Continuous



Continuous



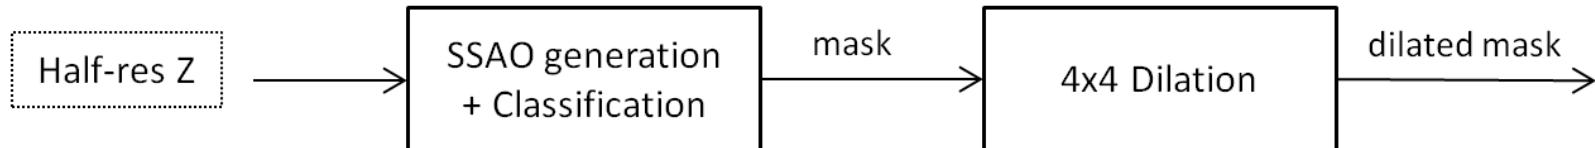
Discontinuous

Two Half-Res Passes

Pass 1: Output SSAO and continuity mask

$$\text{continuityMask} = (\| P_x - P \| < L \&\& \| P_y - P \| < L)$$

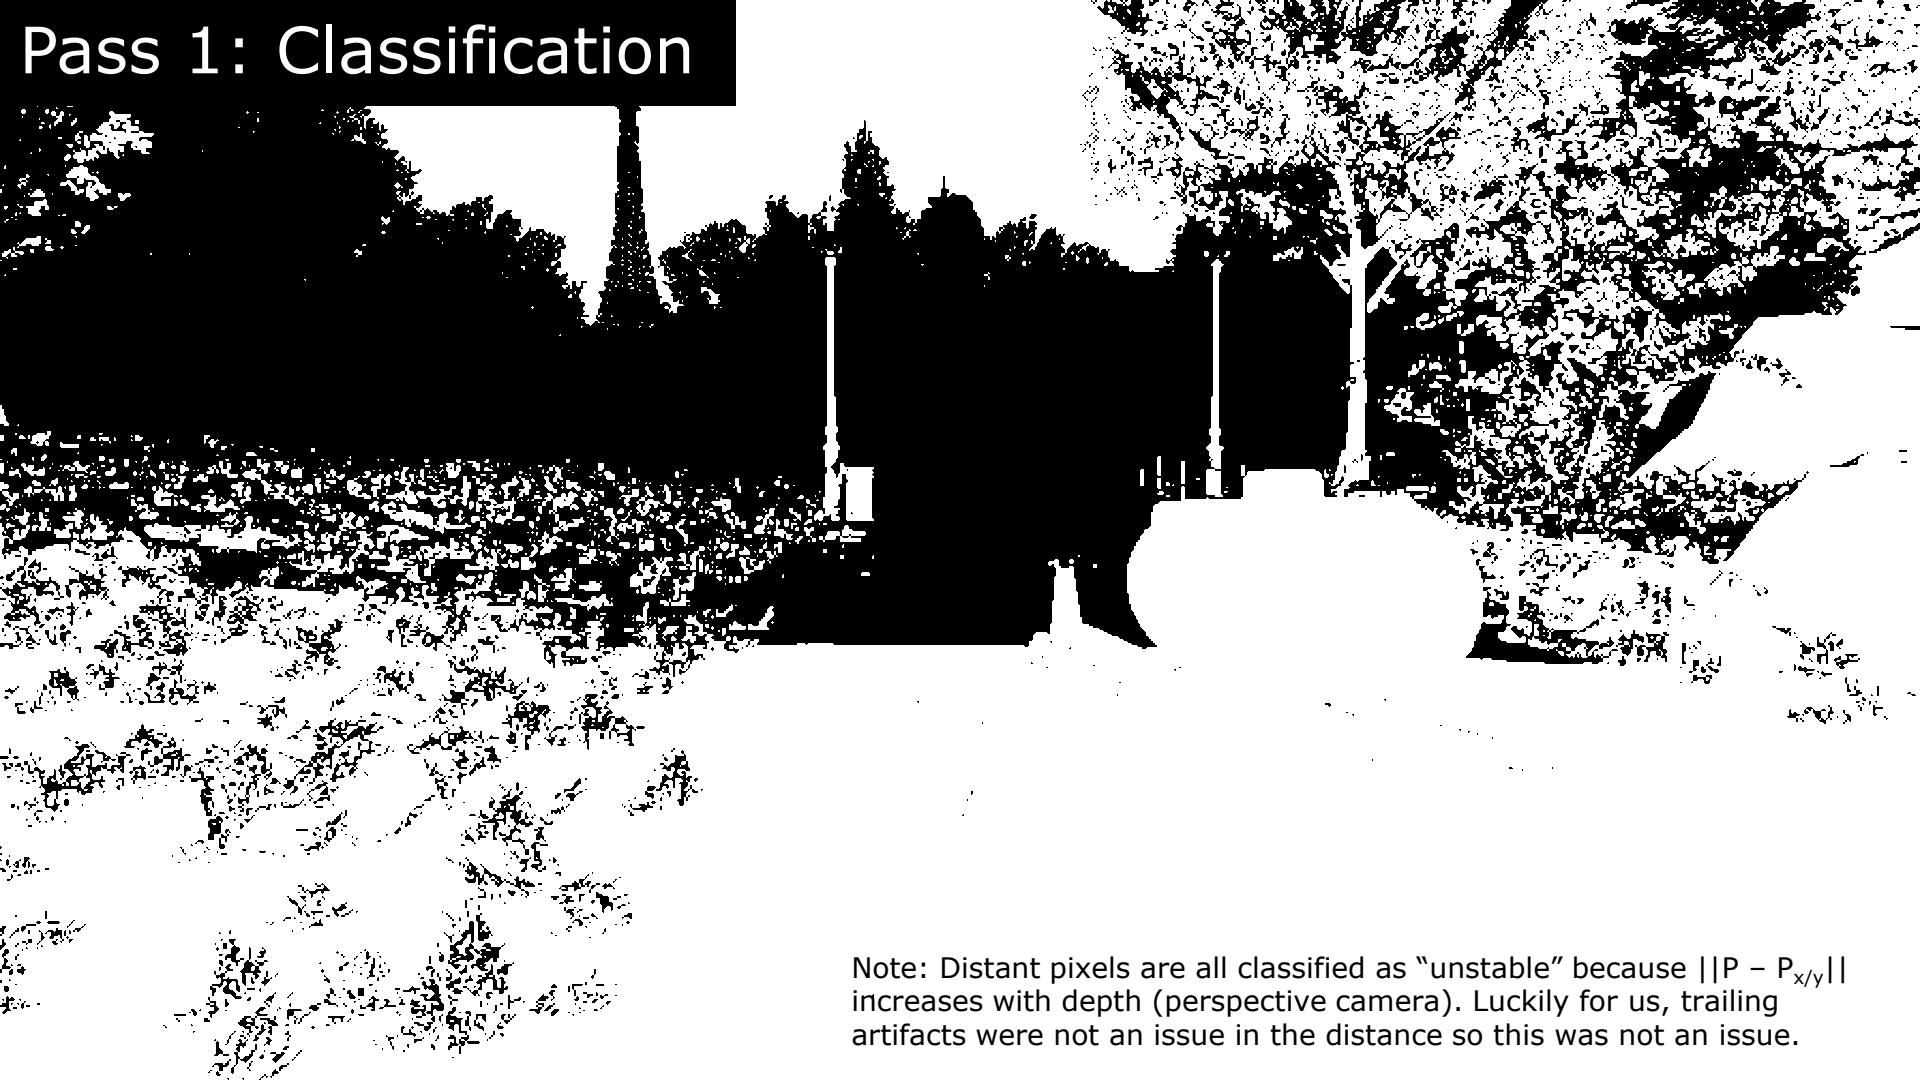
Pass 2: Dilate the discontinuities

$$\text{dilatedMask} = \text{all}_{4x4}(\text{continuityMask})$$


Example Scene



Pass 1: Classification



Note: Distant pixels are all classified as "unstable" because $\|P - P_{x/y}\|$ increases with depth (perspective camera). Luckily for us, trailing artifacts were not an issue in the distance so this was not an issue.

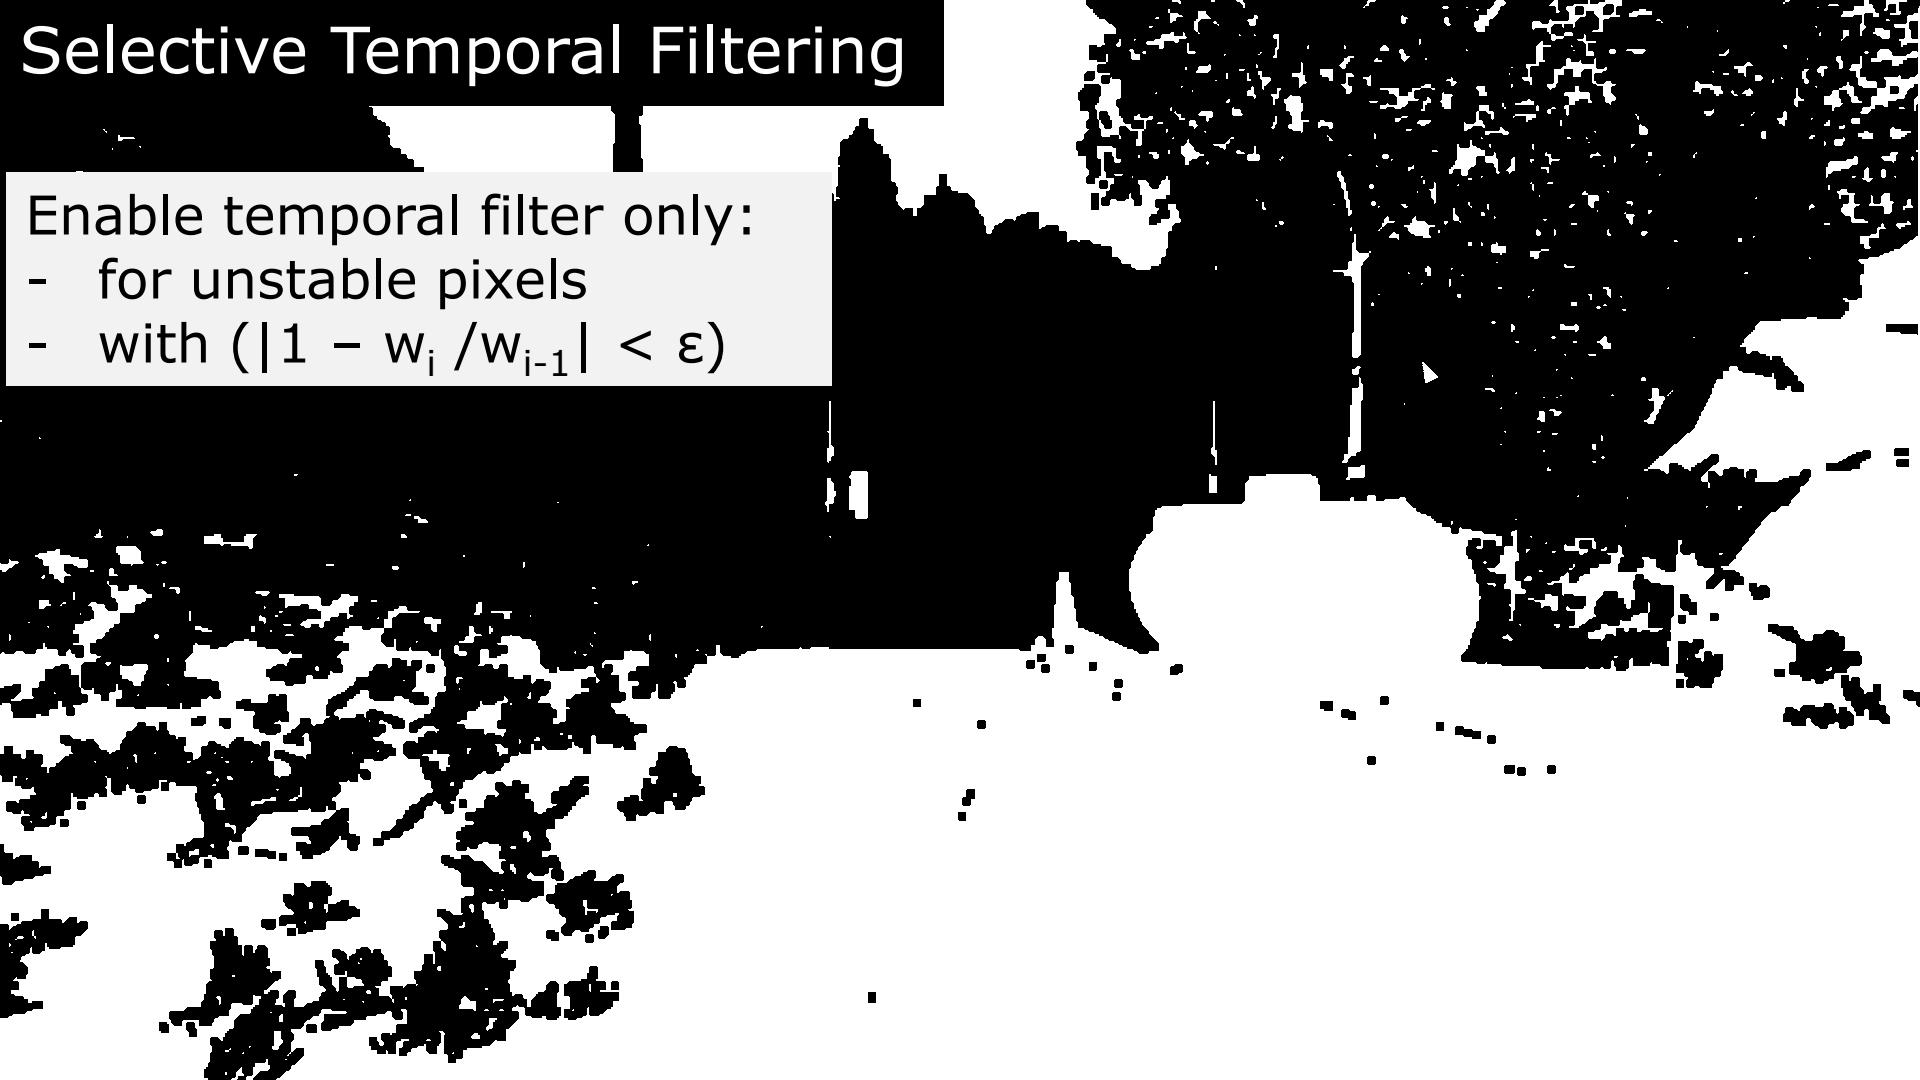
Pass 2: 4x4 Dilation

Pixel classified as unstable if it has at least one discontinuity in its neighborhood

Selective Temporal Filtering

Enable temporal filter only:

- for unstable pixels
- with $(|1 - w_i / w_{i-1}| < \varepsilon)$



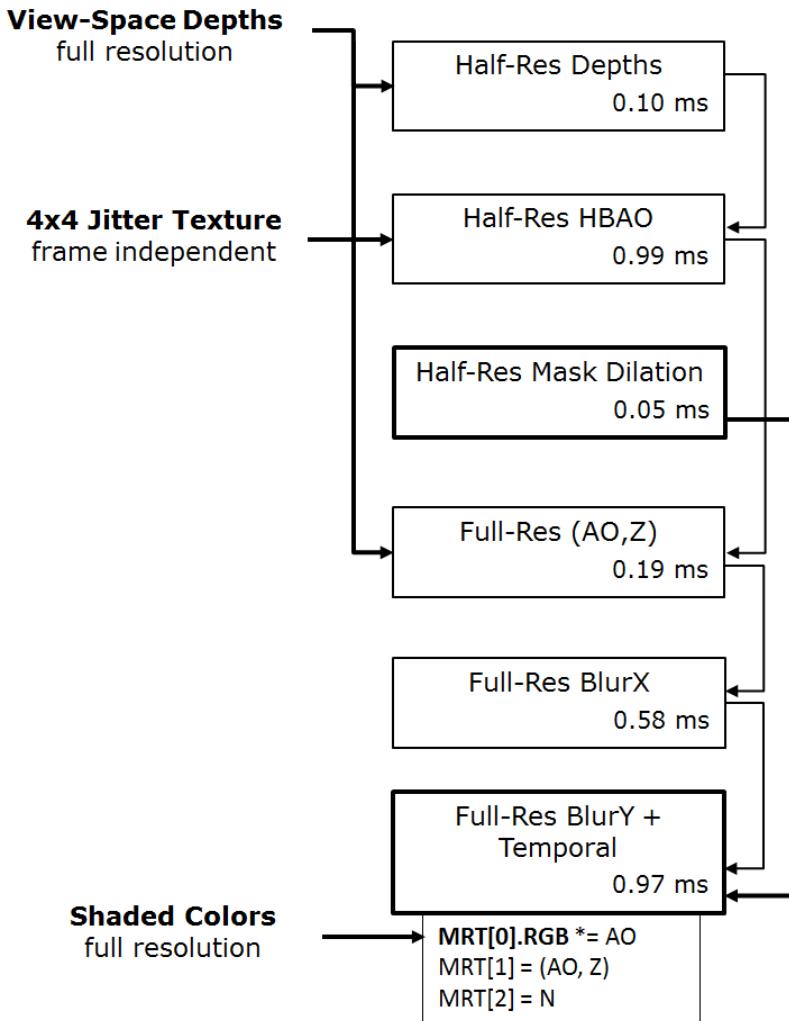
Video: HBAO + Selective Temporal Filtering

BETA



Video: Final Result





Final Pipeline with Selective Temporal Filtering (STF)

STF Performance Hit

[1920x1200 “High”, GTX 560 Ti]

- HBAO total: 2.5 ms -> 2.9 ms
- Frame time (GPU): 25.2 -> 25.6 ms (1.6% performance hit)

STF Parameters

- Reprojection Threshold: For detecting disocclusions ($\varepsilon=10\%$)
- Distance Threshold: For detecting discontinuities ($L=0.1$ meter)
- Dilation Kernel Size (4x4 texels)

History Buffers

- Additional GPU memory required for the history buffers
 - For (AO_i, Z_i, N_i) and (AO_{i-1}, Z_{i-1}, N_i)
 - Full-res, 1xMSAA
- For Multi-GPU Alternate Frame Rendering (AFR) rendering
 - Create one set of buffers per GPU and alternate between them
 - Each AFR GPU has its own buffers & reprojection state
- The history buffers are cleared on first use
 - Clear values: $(AO, Z) = (1.f, 0.f)$ and $N = 0$

SelectiveTemporalFilter(uv_i , AO_i)

```
zi = Fetch(ZBufferi, uvi)
Pi = UnprojectToWorld(uvi, zi)
uvi-1 = ProjectToUV(Viewi-1, Pi)
zi-1 = Fetch(ZBufferi-1, uvi-1)
Pi-1 = UnprojectToWorld(uvi-1, zi-1)
wi = ViewDepth(Viewi-1, Pi)
wi-1 = ViewDepth(Viewi-1, Pi-1)
isStablePixel = Fetch(StabilityMask, uvi)
if ( |1 - wi/wi-1 | < ε && !isStablePixel )
    AOi-1 = Fetch(AOTexturei-1, uvi-1)
    Ni-1 = Fetch(NTexturei-1, uvi-1)
    AOi = (Ni-1 AOi-1 + AOi) / (Ni-1 + 1)
    Ni = min(Ni-1 + 1, Nmax)
else
    Ni = 1
return(AOi, wi, Ni)
```

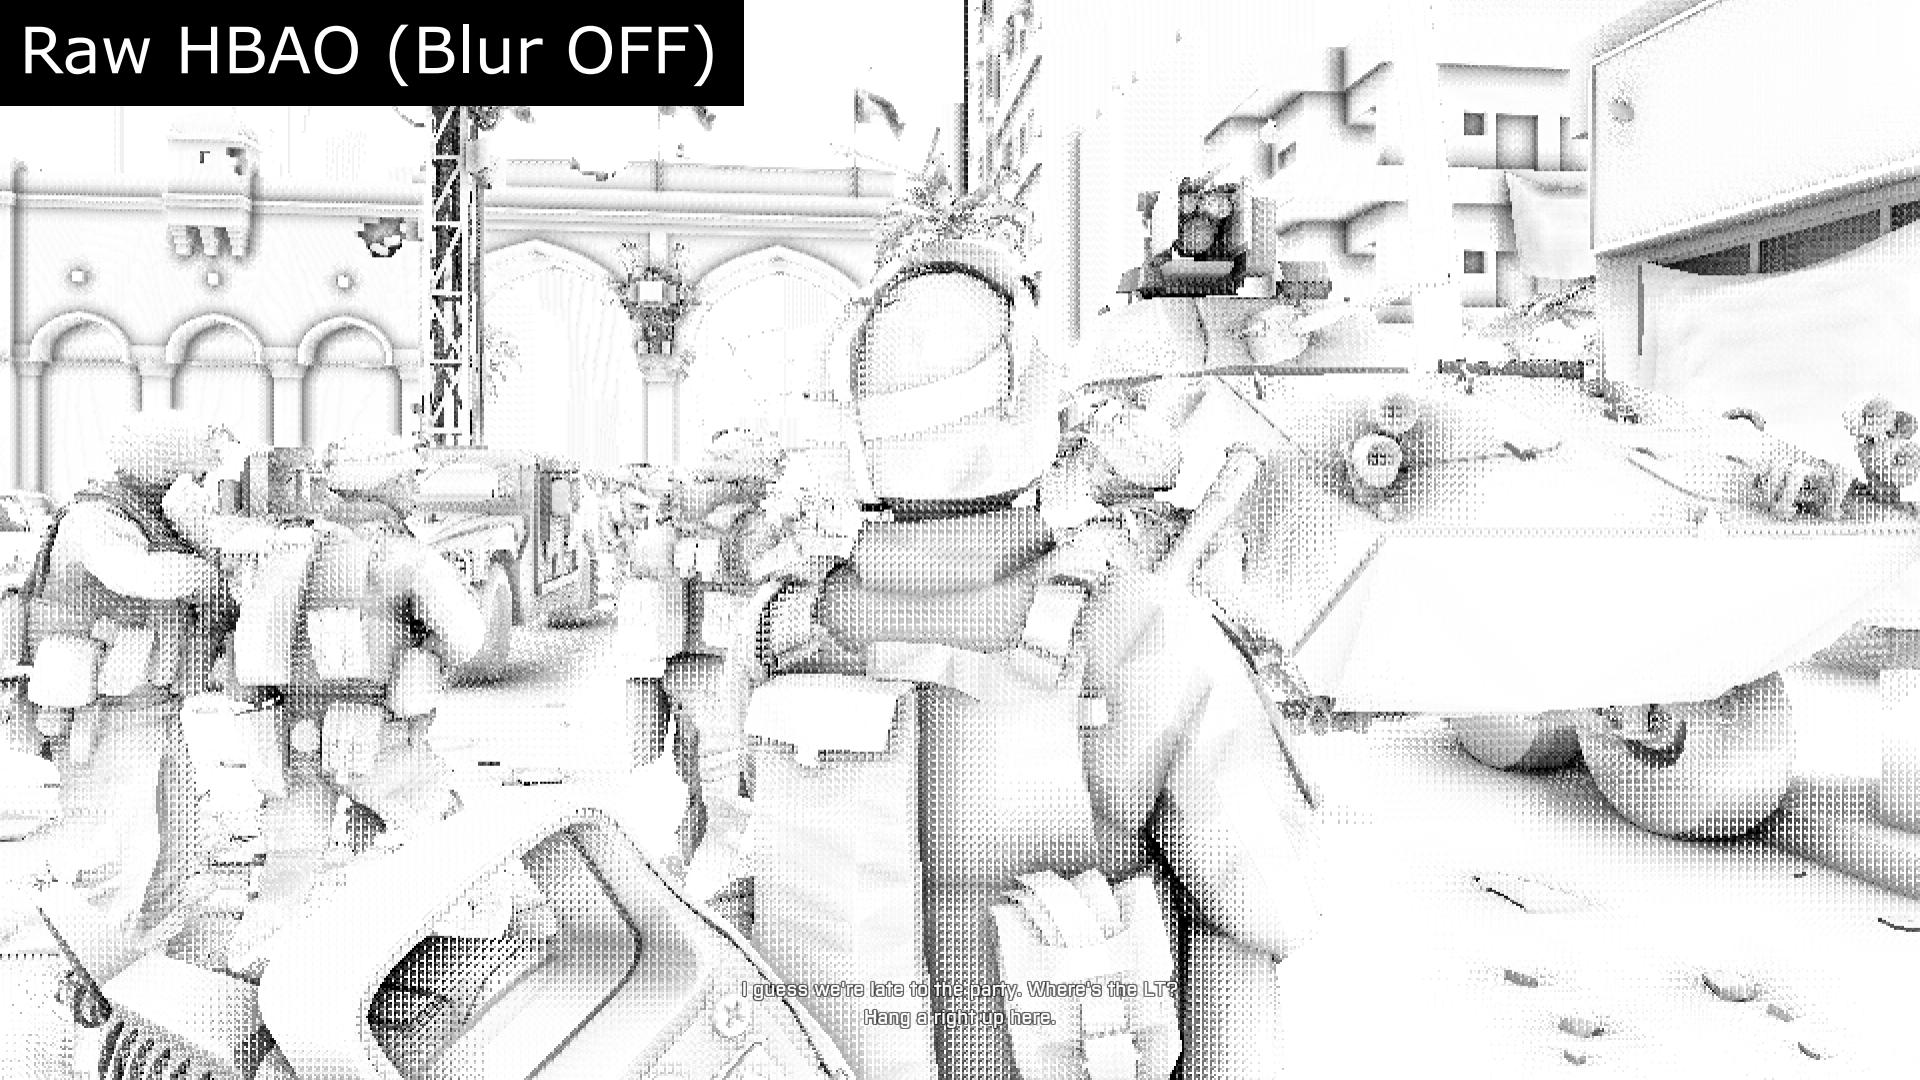
Unoptimized pseudo-code

For fetching z_{i-1} , use **clamp-to-border** to discard out-of-frame data, with $\text{borderZ}=0.f$

For fetching AO_{i-1} , use **bilinear filtering** like in [Nehab et al. 07]

Blur Optimization

Raw HBAO (Blur OFF)



Final HBAO (Blur ON)



I guess we're late to the party. Where's the LT?
Hang a right up here.

Blur Overview

- Full-screen pixel-shader passes
 - BlurX (horizontal)
 - BlurY (vertical)
- BlurX takes as input
 - Half-res AO
 - Full-res linear depth (non-MSAA)

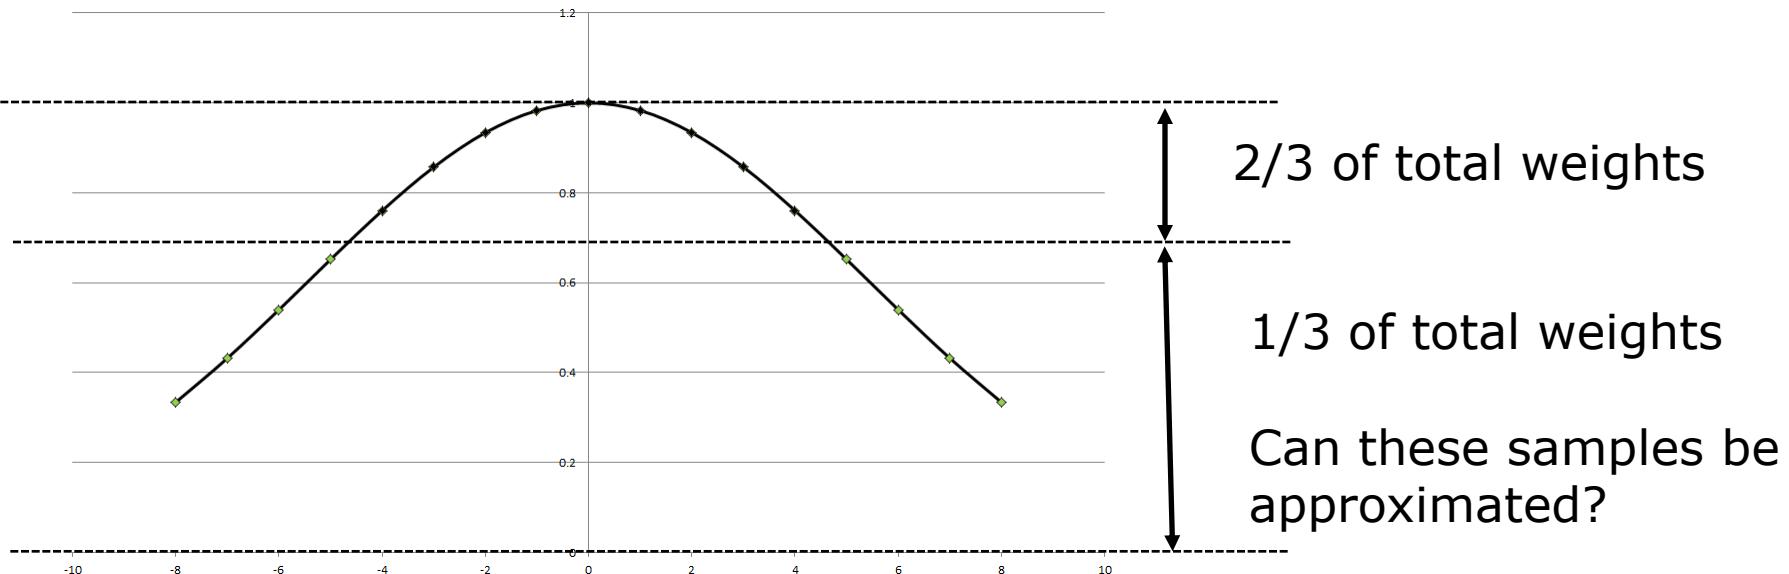
Blur Kernel

- We use 1D Cross-Bilateral Filters (CBF)
- Gaussian blur with depth-dependent weights

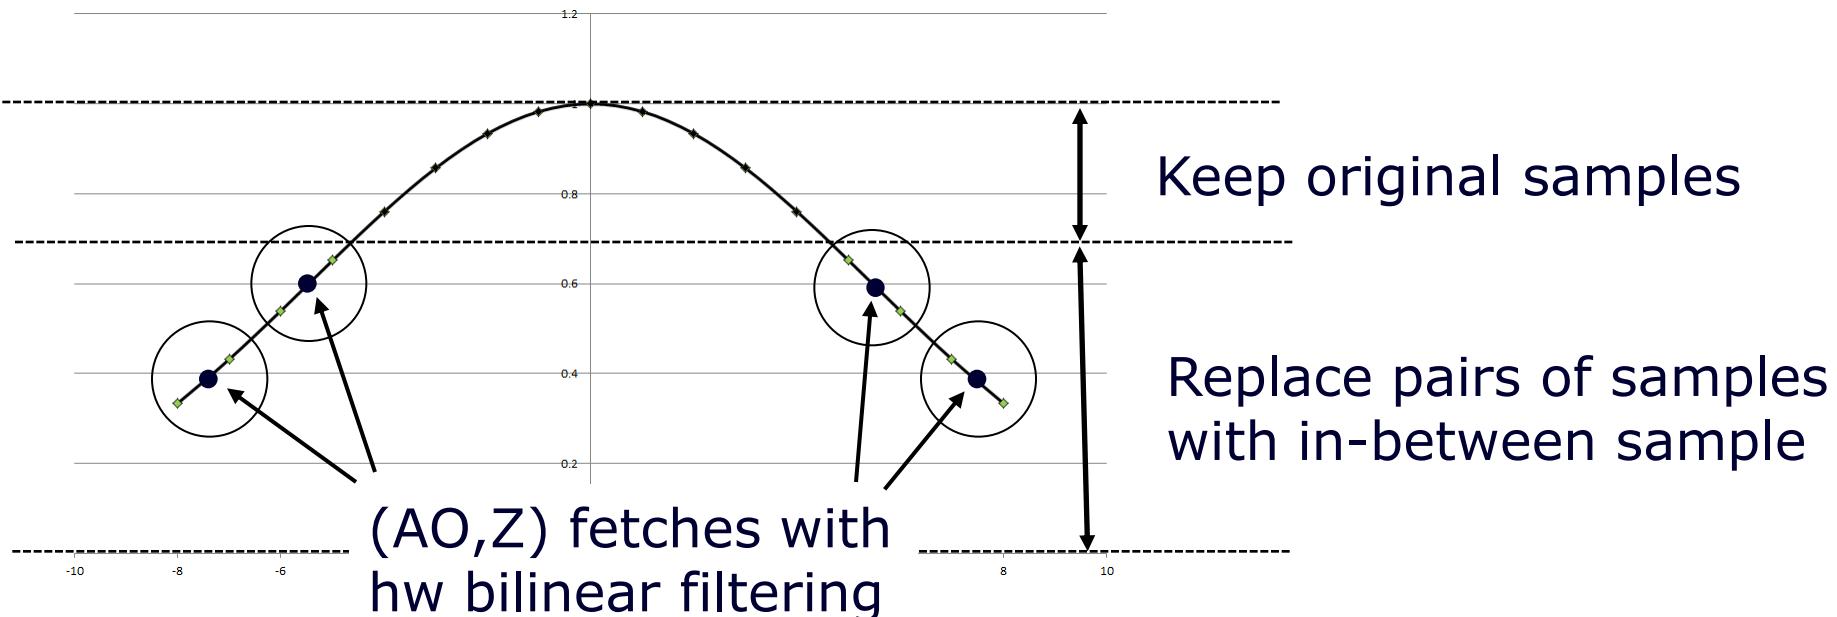
$$\text{Output} = \frac{\text{Sum}[\mathbf{AO}_i w(i, \mathbf{Z}_i, \mathbf{Z}_0), i=-R..R]}{\text{Sum}[w(i, \mathbf{Z}_i, \mathbf{Z}_0), i=-R..R]}$$

[Petschnigg et al. 04]
[Eisemann and Durand 04]
[Kopf et al. 07]
[Bavoil et al. 08]
[McGuire et al. 11]

Blur Opt: Adaptive Sampling



Blur Opt: Adaptive Sampling



Blur Opt: Adaptive Sampling

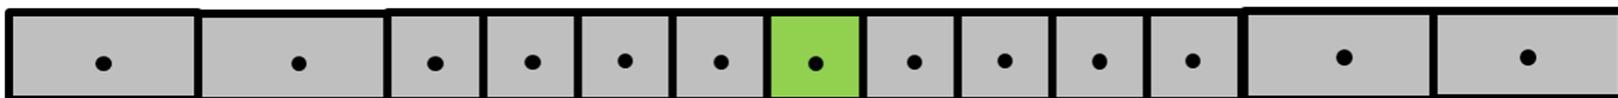
Kernel to be sampled (radius = 8 pixels):



Before: Sample the inner-half of the kernel with POINT filtering and step size = 1:



After: Sample the outer-half of the kernel with LINEAR filtering and step size = 2:



Brute-Force Sampling

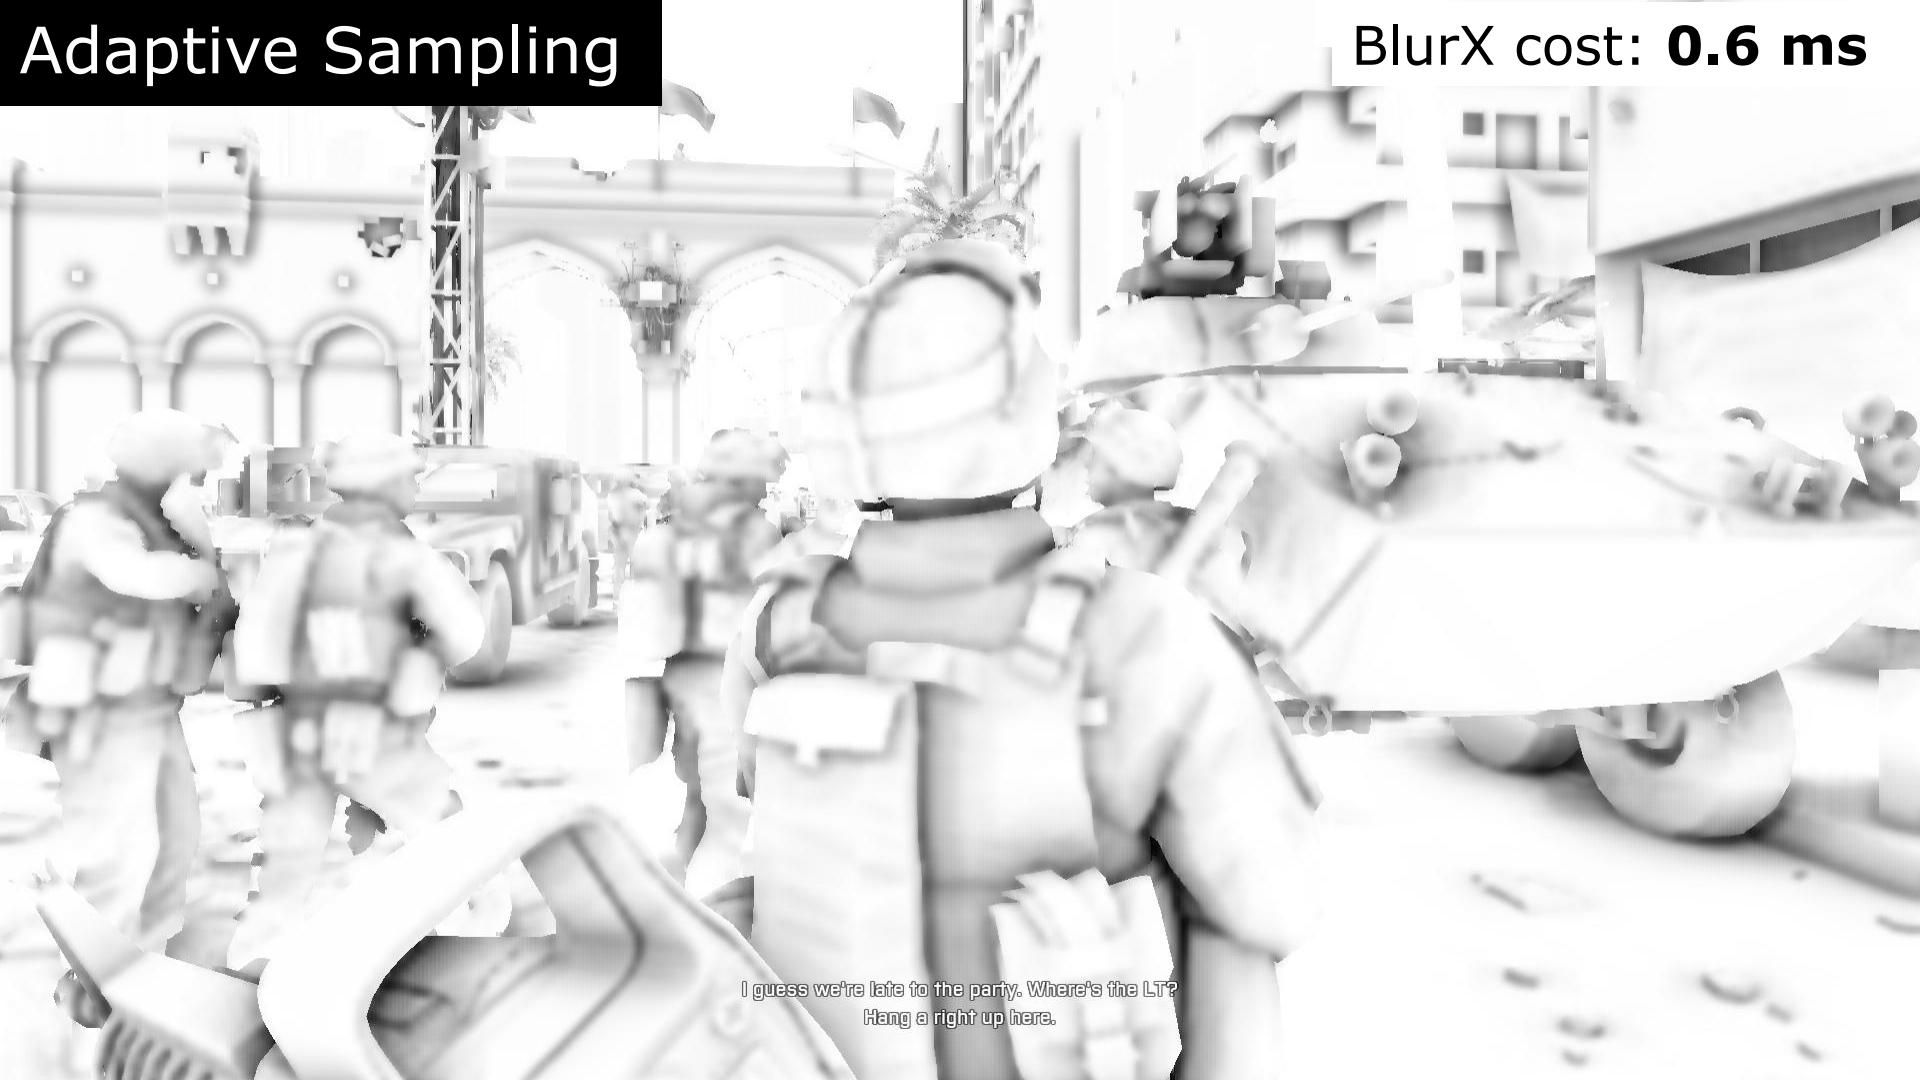
BlurX cost: **0.8 ms**



I guess we're late to the party. Where's the LT?
Hang a right up here.

Adaptive Sampling

BlurX cost: **0.6 ms**



I guess we're late to the party. Where's the LT?
Hang a right up here.

Blur Opt: Speedup

GPU Time	Before	After	Speedup
Pack (AO,Z)	0.18 ms	0.18 ms	0%
BlurX	0.75 ms	0.58 ms	29%
BlurY+STF	1.00 ms	0.95 ms	5% (*)
Blur Total	1.93 ms	1.71 ms	13%

(*) Lower speedup due to the math overhead of
the Selective Temporal Filter (STF)

Blur Radius: 8
Resolution: 1920x1200
GeForce GTX 560 Ti

Summary

Two techniques used in Battlefield 3 / PC

1. A generic solution to fix SSAO flickering with a low perf hit (*) on DX10/11 GPUs
2. An approximate cross-bilateral filter, using a mix of point and bilinear taps

(*) **0.4 ms** in 1920x1200 on GeForce GTX 560 Ti

Questions?

Louis Bavoil
lbavoil@nvidia.com

Johan Andersson
johan.andersson@dice.se



References

- [McGuire et al. 11] McGuire, Osman, Bukowski, Hennessy. The Alchemy Screen-Space Ambient Obscurance Algorithm. Proceedings of ACM SIGGRAPH / Eurographics High-Performance Graphics 2011 (HPG '11).
- [White and Barré-Brisebois 11] White, Barré-Brisebois. More Performance! Five Rendering Ideas from Battlefield 3 and Need For Speed: The Run. Advances in Real-Time Rendering in Games. SIGGRAPH 2011.
- [Mattausch et al. 11] Mattausch, Scherzer, Wimmer. Temporal Screen-Space Ambient Occlusion. In GPU Pro 2. 2011.
- [Loos and Sloan 10] Loos, Sloan. Volumetric Obscurance. ACM Symposium on Interactive 3D Graphics and Games 2010.

References

- [Andersson 10] Andersson. Bending the Graphics Pipeline. Beyond Programmable Shading course, SIGGRAPH 2010.
- [Bavoil and Sainz 09a] Bavoil, Sainz. Image-Space Horizon-Based Ambient Occlusion. In ShaderX7. 2009.
- [Bavoil and Sainz 09b] Bavoil, Sainz. Multi-Layer Dual-Resolution Screen-Space Ambient Occlusion. SIGGRAPH Talk. 2009.
- [Smedberg and Wright 09] Smedberg, Wright. Rendering Techniques in Gears of War 2. GDC 2009.
- [Kajalin 09] Kajalin. Screen Space Ambient Occlusion. In ShaderX7. 2009.
- [Bavoil et al. 08] Bavoil, Sainz, Dimitrov. Image-Space Horizon-Based Ambient Occlusion. SIGGRAPH Talk. 2008.

References

- [Nehab et al. 07] Nehab, Sander, Lawrence, Tatarchuk, Isidoro. Accelerating Real-Time Shading with Reverse Reprojection Caching. In ACM SIGGRAPH/Eurographics Symposium on Graphics Hardware 2007.
- [Kopf et al. 07] Kopf, Cohen, Lischinski, Uyttendaele. 2007. Joint Bilateral Upsampling. In Proceedings of SIGGRAPH 2007.
- [Petschnigg et al. 04] Petschnigg, Szeliski, Agrawala, Cohen, Hoppe. Toyama: Digital photography with flash and no-flash image pairs. In Proceedings of SIGGRAPH 2004.
- [Eisemann and Durand 04] Eisemann, Durand. “Flash Photography Enhancement via Intrinsic Relighting”. In Proceedings of SIGGRAPH 2004.

Bonus Slides

HLSL: Adaptive Sampling

```
float r = 1;  
// Inner half of the kernel: step size = 1 and POINT filtering  
[unroll] for (; r <= KERNEL_RADIUS/2; r += 1)  
{  
    float2 uv = r * deltaUV + uv0;  
    float2 AOZ = mainTexture.Sample(pointClampSampler, uv).xy;  
    processSample(AOZ, r, centerDepth, totalAO, totalW);  
}  
// Outer half of the kernel: step size = 2 and LINEAR filtering  
[unroll] for (; r <= KERNEL_RADIUS; r += 2)  
{  
    float2 uv = (r + 0.5) * deltaUV + uv0;  
    float2 AOZ = mainTexture.Sample(linearClampSampler, uv).xy;  
    processSample(AOZ, r, centerDepth, totalAO, totalW);  
}
```

HLSL: Cross-Bilateral Weights

```
// d and d0 = linear depths
float crossBilateralWeight(float r, float d, float d0)
{
    // precompiled by fxc
    const float BlurSigma = ((float)KERNEL_RADIUS+1.0f) * 0.5f;
    const float BlurFalloff = 1.f / (2.0f*BlurSigma*BlurSigma);

    // assuming that d and d0 are pre-scaled linear depths
    float dz = d0 - d;
    return exp2(-r*r*BlurFalloff - dz*dz);
}
```